

APPENDIX: DETAILED REVENUE INITIATIVES

Rank	Initiative	Solution Summary	Customer Result	Google Cloud pull-through	Estimated annual revenue (USD M)	Weekly active users added
1	Cloud-credit swap	The field organisation attaches Gemini Code Assist Enterprise seats to any new or expanded committed-use agreement. It activates those seats immediately through existing license-assignment tools in the Google Cloud console.	Every developer in the purchasing company receives the assistant without a separate procurement action. It begins using it inside Cloud Workstations or the Google VS Code extension on the first working day after contract signature.	The larger commitment increases contracted Google Cloud consumption for compute, storage, and networking and links renewal decisions to the assistant's ongoing value.	500	600,000
2	Vertex fine-tune lane	Customers use the current Vertex AI fine-tuning service to train a private Gemini model on their internal repositories and set that private endpoint in the Code Assist configuration panel.	Suggestions follow the organisation's naming conventions and frameworks, so review cycles shorten and merge rates improve.	Every training job bills Vertex AI Training hours, and every suggestion against the private model bills Vertex AI inference tokens, producing a continuous revenue stream.	145	120,000

3	Vertex backend with Cursor	Cursor users select "Google models" in the existing settings menu; the IDE routes each prompt to Gemini on Vertex AI under the customer's Google Cloud project.	Developer teams keep their preferred Cursor workflow while gaining access to Gemini quality without installing extra infrastructure.	All prompt traffic is metered as Vertex inference tokens, adding predictable consumption revenue without new features or operational overhead.	18	60,000
4	GKE auto-YAML	Inside Cloud Workstations, the assistant generates valid Kubernetes manifests; developers copy the files into repositories and deploy them with today's Cloud Deploy pipelines.	Teams launch services on Google Kubernetes Engine without writing YAML by hand, reducing setup errors and saving build minutes.	Each new cluster incurs the standard GKE control-plane fee, and each Cloud Deploy run records pipeline charges, producing incremental platform revenue tied to assistant usage.	5	30,000
5	Marketplace rev-share for plug-ins	Third-party vendors list Gemini Code Assist extensions in Cloud Marketplace and receive eighty-five percent of sales while Google keeps the standard share.	Customers can purchase domain-specific helpers, such as FPGA snippets or industry templates, through a single billing channel, making add-on adoption easier.	Marketplace transactions count toward the commitment and create a new micro-ecosystem that drives additional Cloud billing lines.	1	8,000

6	Reference architecture on GKE with Cursor	Google publishes a Terraform script that spins up GKE, Cloud SQL, and a pre-configured Cursor IDE plug-in that already points to Gemini on Vertex AI.	Enterprise pilots can finish setting up in one hour instead of several days and measure productivity gains in the same week.	Each deployment produces GKE node hours, Cloud SQL storage, and Vertex token spending, all tracked under the customer's existing account.	1	6,000
7	GCP Marketplace listing for Cursor	Cursor's SaaS entry appears in Cloud Marketplace with a "Works with Gemini" badge, allowing buyers to draw the subscription from their committed-use credits.	Procurement teams complete the purchase in minutes and align billing with their existing Google Cloud invoices.	Spends routed through the marketplace count toward the customer's spending commitment and Google Cloud's normal revenue share.	1	4,000
8	Revenue-share seat bundle	Google field sellers offer an "AI Dev Pack" that bundles one Cursor seat and one Code Assist Enterprise seat on a single order form with a predefined revenue split.	Large enterprises receive unified pricing and a single point of contact, removing a hurdle that slows multi-vendor deals.	The bundle adds Code Assist seats and books Cursor revenue through Marketplace, lifting license revenue and related cloud usage.	<1	3,000

9	Event co-marketing	Google and Cursor run a quarterly online coding challenge in which human teams use both tools to complete real tasks. Winners receive Google Cloud credits and free Cursor time.	Participants see direct speed gains, share experiences in their organisations, and often adopt one or both tools after the event.	Redeemed credits convert into paid workloads for Vertex AI, GKE, or Cloud Run when trial projects move to production.	<1	2,000
10	Shared ROI assessment	Google solution architects and Cursor engineers run a four-week proof-of-value that records baselines in Jira and shows time saved once Gemini-backed Cursor is enabled.	The customer receives a quantified business case that supports an enterprise-wide roll-out decision.	Successful studies typically precede a cloud-credit swap or Marketplace purchase, so the programme feeds higher-revenue plays without new product work.	<1	1,000

Rank	Sales Play	Solution Summary	Customer Result	Google Cloud pull-through	Annual revenue mid-case (USD M)	Weekly active users added
10	Data-governance add-on (Cursor)	Run Cursor IDE in Anthos or GKE and route prompts to a private Vertex endpoint through Private Service Connect.	Regulated teams keep source code and model traffic inside their VPC and still use Gemini suggestions.	Anthos or GKE license plus Vertex inference tokens recorded under the customer project.	80	45 000

11	BigQuery SQL helper	Use the existing BigQuery Console plug-in so Gemini writes and tunes SQL queries.	Analysts finish queries faster and run more iterations.	Extra queries consume BigQuery slots and storage.	36	200 000
12	Cloud Run cost guard	In Cloud Workstations, Gemini rewrites handlers to reduce cold starts and redeploys through Cloud Deploy.	Services start faster and cost less to run.	More Cloud Run services monitoring metrics, plus Vertex tokens per suggestion.	12	40 000
13	Secure patch loop	Security Command Center sends findings to Code Assist, which drafts pull requests for fixes.	Vulnerabilities are patched sooner without manual edits.	Upsells SCC and bills Vertex tokens for each generated patch.	9	25 000
14	Apigee API boilerplate	Gemini reads an OpenAPI spec stored in Apigee, generates client libraries and unit tests, and commits them to the repo.	API teams publish SDKs quickly and with consistent tests.	Client calls raise Apigee traffic; each generation uses Vertex tokens.	5	15 000
15	Workspace doc-to-repo	A Google Docs add-on calls Gemini to scaffold a new repo in Cloud Source Repositories and create a starter Cloud Build pipeline.	Small teams create a compliant repo structure in minutes.	New repos, Cloud Build jobs, and Artifact Registry pulls record platform usage.	2	12 000

16	Anthos edge bundle	Ship an on-prem Anthos rack with an offline Gemini model that matches today's Anthos VM spec.	Developers in disconnected sites keep using the assistant without sending data to the public cloud.	Annual Anthos subscription revenue and optional support contract.	1	4 000
17	Usage-based discount trigger (Cursor)	When the Vertex token spend driven by Cursor crosses a preset threshold, the account unlocks extra Cursor seats at no charge.	Heavy users gain more seats when they increase usage rather than pay per seat.	Incentive drives higher Vertex inference volume on existing infrastructure.	1	5 000
18	Cloud Run cost guard usage trigger	A scale metric in Cloud Monitoring triggers automatic handler optimisation when cold-start latency exceeds a limit.	Services remain cost-efficient without manual tuning.	Each optimisation event records Vertex tokens and reinforces Cloud Run stickiness.	1	5 000
19	Support escalation pact	Google and Cursor operate a shared Slack channel and follow a joint run-book for incidents.	Customers get one escalation path instead of juggling vendors.	Pact helps retain existing seats and workloads, though it adds little new spend.	0	0